

## Classification using Logistic Regression

Marquez, Keith Leigh Zhen R.

Objective(s):

This activity aims to demonstrate how to apply simple linear regression analysis to solve regression problem

```
pip install ucimlrepo

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)

from ucimlrepo import fetch_ucirepo

# fetch dataset
cervical_cancer_risk_factors = fetch_ucirepo(id=383)

# data (as pandas dataframes)
X = cervical_cancer_risk_factors.data.features
y = cervical_cancer_risk_factors.data.targets

# metadata
print(cervical_cancer_risk_factors.metadata)

# variable information
print(cervical_cancer_risk_factors.variables)

 14      STDs:cervical condylomatosis  Feature  Continuous      None
 15      STDs:vaginal condylomatosis  Feature  Continuous      None
 16  STDs:vulvo-perineal condylomatosis  Feature  Continuous      None
 17                  STDs:syphilis  Feature  Continuous      None
 18  STDs:pelvic inflammatory disease  Feature  Continuous      None
 19          STDs:genital herpes  Feature  Continuous      None
 20      STDs:molluscum contagiosum  Feature  Continuous      None
 21                  STDs:AIDS  Feature  Continuous      None
 22                  STDs:HIV  Feature  Continuous      None
 23          STDs:Hepatitis B  Feature  Continuous      None
 24          STDs:HPV  Feature  Continuous      None
 25      STDs: Number of diagnosis  Feature  Integer      None
 26  STDs: Time since first diagnosis  Feature  Continuous      None
 27  STDs: Time since last diagnosis  Feature  Continuous      None
 28          Dx:Cancer  Feature  Integer      None
 29          Dx:CIN  Feature  Integer      None
 30          Dx:HPV  Feature  Integer      None
 31          Dx  Feature  Integer      None
 32          Hinselmann  Feature  Integer      None
 33          Schiller  Feature  Integer      None
 34          Cytology  Feature  Integer      None
 35          Biopsy  Feature  Integer      None

  description  units  missing_values
 0        None    None         no
 1        None    None        yes
 2        None    None        yes
 3        None    None        yes
 4        None    None        yes
 5        None    None        yes
 6        None    None        yes
 7        None    None        yes
```

```

25    None  None      no
26    None  None     yes
27    None  None     yes
28    None  None      no
29    None  None      no
30    None  None      no
31    None  None      no
32    None  None      no

```

## ▼ Data Wrangling

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
cc = pd.concat([X,y], axis=1)
cc

```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs: Time since first diagnosis	S: diagnosis
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	
2	34	1.0	NaN	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	NaN	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	NaN	
...	...	...	...	...	...	...	...	...	...	...	...	...	
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	NaN	
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	NaN	
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	NaN	
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	NaN	
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	NaN	

858 rows × 36 columns

cc.dtypes

Age	int64
Number of sexual partners	float64
First sexual intercourse	float64
Num of pregnancies	float64
Smokes	float64
Smokes (years)	float64
Smokes (packs/year)	float64
Hormonal Contraceptives	float64
Hormonal Contraceptives (years)	float64
IUD	float64
IUD (years)	float64
STDs	float64
STDs (number)	float64
STDs:condylomatisis	float64
STDs:cervical condylomatisis	float64
STDs:vaginal condylomatisis	float64
STDs:vulvo-perineal condylomatisis	float64
STDs:syphilis	float64
STDs:pelvic inflammatory disease	float64
STDs:genital herpes	float64
STDs:molluscum contagiosum	float64
STDs:AIDS	float64
STDs:HIV	float64
STDs:Hepatitis B	float64
STDs:HPV	float64
STDs: Number of diagnosis	int64
STDs: Time since first diagnosis	float64
STDs: Time since last diagnosis	float64

```
Dx:Cancer          int64
Dx:CIN            int64
Dx:HPV            int64
Dx                int64
Hinselmann        int64
Schiller          int64
Citology          int64
Biopsy            int64
dtype: object
```

```
cc_null=cc.isnull().sum()
cc_null

Age                  0
Number of sexual partners    26
First sexual intercourse     7
Num of pregnancies          56
Smokes                 13
Smokes (years)           13
Smokes (packs/year)       13
Hormonal Contraceptives    108
Hormonal Contraceptives (years) 108
IUD                   117
IUD (years)             117
STDs                  105
STDs (number)           105
STDs:condylomatosis      105
STDs:cervical condylomatosis 105
STDs:vaginal condylomatosis 105
STDs:vulvo-perineal condylomatosis 105
STDs:syphilis            105
STDs:pelvic inflammatory disease 105
STDs:genital herpes       105
STDs:molluscum contagiosum 105
STDs:AIDS                105
STDs:HIV                 105
STDs:Hepatitis B          105
STDs:HPV                 105
STDs: Number of diagnosis 0
STDs: Time since first diagnosis 787
STDs: Time since last diagnosis 787
Dx:Cancer               0
Dx:CIN                 0
Dx:HPV                 0
Dx                0
Hinselmann            0
Schiller              0
Citology              0
Biopsy                0
dtype: int64
```

```
for column in cc.columns:
    cc[column] = cc[column].fillna(cc[column].mode()[0])
cc.isnull().sum()
```

```
Age                  0
Number of sexual partners    0
First sexual intercourse     0
Num of pregnancies          0
Smokes                 0
Smokes (years)           0
Smokes (packs/year)       0
Hormonal Contraceptives    0
Hormonal Contraceptives (years) 0
IUD                   0
IUD (years)             0
STDs                  0
STDs (number)           0
STDs:condylomatosis      0
STDs:cervical condylomatosis 0
STDs:vaginal condylomatosis 0
STDs:vulvo-perineal condylomatosis 0
STDs:syphilis            0
STDs:pelvic inflammatory disease 0
STDs:genital herpes       0
STDs:molluscum contagiosum 0
STDs:AIDS                0
STDs:HIV                 0
STDs:Hepatitis B          0
STDs:HPV                 0
STDs: Number of diagnosis 0
```

```

STDs: Time since first diagnosis      0
STDs: Time since last diagnosis     0
Dx:Cancer                           0
Dx:CIN                             0
Dx:HPV                            0
Dx                                0
Hinselmann                         0
Schiller                           0
Cytology                           0
Biopsy                            0
dtype: int64

```

cc

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs: Time since first diagnosis	S:
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0	
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0	
2	34	1.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0	
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	1.0	
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	1.0	
...	...	...	...	...	...	...	...	...	...	...	...	...	
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0	
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	1.0	
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	1.0	
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	1.0	
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	1.0	

858 rows × 36 columns

## ▼ Remove duplicates

```

# Check for duplicates
duplicate_rows = cc.duplicated()

# Count the number of duplicate rows
num_duplicates = duplicate_rows.sum()
print("Number of duplicate rows:", num_duplicates)

```

Number of duplicate rows: 28

```

cc.drop_duplicates(inplace=True)
cc

```

Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs:	S:
											Time since first diagnosis	diagnosis
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0
2	34	1.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	...	1.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	...	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	...	1.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	...	1.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	...	1.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	...	1.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	...	1.0

830 rows × 36 columns


cc.describe()

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	
count	830.000000	830.000000	830.000000	830.000000	830.000000	830.000000	830.000000	830.000000	830.000000	830.000000	
mean	27.075904	2.540964	17.010843	2.222892	0.148193	1.241765	0.461333	0.696386	2.031005	0.100000	
std	8.479609	1.656987	2.817510	1.445331	0.355505	4.122522	2.245822	0.460095	3.639632	0.300181	
min	13.000000	1.000000	10.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	21.000000	2.000000	15.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	26.000000	2.000000	17.000000	2.000000	0.000000	0.000000	0.000000	1.000000	0.250000	0.000000	
75%	32.000000	3.000000	18.000000	3.000000	0.000000	0.000000	0.000000	1.000000	3.000000	0.000000	
max	84.000000	28.000000	32.000000	11.000000	1.000000	37.000000	37.000000	1.000000	30.000000	1.000000	

8 rows × 36 columns


## EDA

```
plt.figure(figsize=(8, 6))
sns.boxplot(x='Dx:Cancer', y='Age', data=cc)
plt.xlabel('Cancer Diagnosis')
plt.ylabel('Age')
plt.title('Age by Cancer Diagnosis')
plt.show()
```




This compares the age distribution of individuals with and without a cancer diagnosis. Both groups have similar median ages around 40, but the non-diagnosed group has a wider range of ages due to outliers.

```
plt.scatter(cc['Age'], cc['Num of pregnancies'], alpha=0.5)
plt.xlabel('Age')
plt.ylabel('Num of pregnancies')
plt.title('Scatter plot between Age and Num of pregnancies')
plt.show()
```




Most pregnancies occur between ages 20 and 50, with fewer occurrences in individuals below 20 or above 50. Age significantly influences pregnancy likelihood, with peak reproductive years around 20 to 50.

```
plt.figure(figsize=(8, 6))
sns.countplot(x='Smokes', hue='Dx:Cancer', data=cc)
plt.xlabel('Smoking Status')
plt.ylabel('Count')
plt.title('Bar Plot of Smoking Status by Cancer Diagnosis')
plt.show()
```



Non-smokers (smoking status 0.0) are less likely to be diagnosed with cancer, while smokers (smoking status 1.0) have a higher proportion of cancer diagnoses. The data suggests a potential link between smoking and cancer risk.

```
# Pairplot for selected numerical columns
sns.pairplot(cc[['Age', 'Number of sexual partners', 'Num of pregnancies', 'Dx:Cancer']], hue='Dx:Cancer')
plt.show()
```



This indicates that most individuals are young adults with fewer sexual partners and pregnancies. A small proportion has been diagnosed with cancer, but there is no clear pattern suggesting a direct association between these factors and cancer risk from the graphs alone.

## Simple Linear Regression and Demonstration of Logistic Regression

### Split data into separate training and test set

```
X = cc.drop('Biopsy', axis=1)
y = cc['Biopsy']

from sklearn.model_selection import train_test_split

# Split x and y into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Check the shape of train and test sets
print("Shape of X_train:", X_train.shape)
print("Shape of X_test:", X_test.shape)

Shape of X_train: (664, 35)
Shape of X_test: (166, 35)
```

```
X_train.dtypes
```

Age	int64
Number of sexual partners	float64
First sexual intercourse	float64
Num of pregnancies	float64
Smokes	float64

```

Smokes (years)           float64
Smokes (packs/year)      float64
Hormonal Contraceptives   float64
Hormonal Contraceptives (years) float64
IUD                      float64
IUD (years)               float64
STDs                     float64
STDs (number)             float64
STDs:condylomatosis       float64
STDs:cervical condylomatosis float64
STDs:vaginal condylomatosis float64
STDs:vulvo-perineal condylomatosis float64
STDs:syphilis              float64
STDs:pelvic inflammatory disease float64
STDs:genital herpes        float64
STDs:molluscum contagiosum float64
STDs:AIDS                  float64
STDs:HIV                   float64
STDs:Hepatitis B            float64
STDs:HPV                   float64
STDs: Number of diagnosis    int64
STDs: Time since first diagnosis float64
STDs: Time since last diagnosis float64
Dx:Cancer                  int64
Dx:CIN                      int64
Dx:HPV                      int64
Dx                          int64
Hinselmann                 int64
Schiller                    int64
Cytology                    int64
dtype: object

```

```

#display numerical variables
numerical = [col for col in X_train.columns if X_train[col].dtypes != "O"]
numerical

['Age',
 'Number of sexual partners',
 'First sexual intercourse',
 'Num of pregnancies',
 'Smokes',
 'Smokes (years)',
 'Smokes (packs/year)',
 'Hormonal Contraceptives',
 'Hormonal Contraceptives (years)',
 'IUD',
 'IUD (years)',
 'STDs',
 'STDs (number)',
 'STDs:condylomatosis',
 'STDs:cervical condylomatosis',
 'STDs:vaginal condylomatosis',
 'STDs:vulvo-perineal condylomatosis',
 'STDs:syphilis',
 'STDs:pelvic inflammatory disease',
 'STDs:genital herpes',
 'STDs:molluscum contagiosum',
 'STDs:AIDS',
 'STDs:HIV',
 'STDs:Hepatitis B',
 'STDs:HPV',
 'STDs: Number of diagnosis',
 'STDs: Time since first diagnosis',
 'STDs: Time since last diagnosis',
 'Dx:Cancer',
 'Dx:CIN',
 'Dx:HPV',
 'Dx',
 'Hinselmann',
 'Schiller',
 'Cytology']

```

## Engineering outliers in numerical variables

```

def max_value(df, variable, top):
    return np.where(df[variable] > top, top, df[variable])

numerical_columns = [
    'Age', 'Number of sexual partners', 'First sexual intercourse', 'Num of pregnancies',
    'Smokes', 'Smokes (years)', 'Smokes (packs/year)', 'Hormonal Contraceptives',
    'Hormonal Contraceptives (years)', 'IUD', 'IUD (years)', 'STDs', 'STDs (number)',
    'STDs:condylomatosis', 'STDs:cervical condylomatosis', 'STDs:vaginal condylomatosis',
    'STDs:vulvo-perineal condylomatosis', 'STDs:syphilis', 'STDs:pelvic inflammatory disease',
    'STDs:genital herpes', 'STDs:molluscum contagiosum', 'STDs:AIDS', 'STDs:HIV',
    'STDs:Hepatitis B', 'STDs:HPV', 'STDs: Number of diagnosis', 'STDs: Time since first diagnosis',
    'STDs: Time since last diagnosis', 'Dx:Cancer', 'Dx:CIN', 'Dx:HPV', 'Dx', 'Hinselmann',
    'Schiller', 'Citology'
]

# Apply max_value function to each numerical column
for df in [X_train, X_test]:
    df['Age'] = max_value(df, 'Age', 100) # Assuming 100 is the maximum allowable age
    df['Number of sexual partners'] = max_value(df, 'Number of sexual partners', 10) # Example threshold
    # Repeat this process for other numerical columns

```

Maximum values after transformation:

```
# Check the maximum values after transformation
print("X_train:", X_train[numerical_columns].max())
```

X_train: Age	70.0
Number of sexual partners	10.0
First sexual intercourse	32.0
Num of pregnancies	8.0
Smokes	1.0
Smokes (years)	37.0
Smokes (packs/year)	37.0
Hormonal Contraceptives	1.0
Hormonal Contraceptives (years)	30.0
IUD	1.0
IUD (years)	19.0
STDs	1.0
STDs (number)	4.0
STDs:condylomatosis	1.0
STDs:cervical condylomatosis	0.0
STDs:vaginal condylomatosis	1.0
STDs:vulvo-perineal condylomatosis	1.0
STDs:syphilis	1.0
STDs:pelvic inflammatory disease	1.0
STDs:genital herpes	1.0
STDs:molluscum contagiosum	1.0
STDs:AIDS	0.0
STDs:HIV	1.0
STDs:Hepatitis B	1.0
STDs:HPV	1.0
STDs: Number of diagnosis	3.0
STDs: Time since first diagnosis	22.0
STDs: Time since last diagnosis	22.0
Dx:Cancer	1.0
Dx:CIN	1.0
Dx:HPV	1.0
Dx	1.0
Hinselmann	1.0
Schiller	1.0
Citology	1.0
dtype: float64	

```
print("X_test:", X_test[numerical_columns].max())
```

X_test: Age	84.0
Number of sexual partners	10.0
First sexual intercourse	28.0
Num of pregnancies	11.0
Smokes	1.0
Smokes (years)	24.0
Smokes (packs/year)	21.0
Hormonal Contraceptives	1.0
Hormonal Contraceptives (years)	22.0
IUD	1.0
IUD (years)	10.0
STDs	1.0
STDs (number)	2.0

```

STDs:condylomatosis      1.0
STDs:cervical condylomatosis 0.0
STDs:vaginal condylomatosis 0.0
STDs:vulvo-perineal condylomatosis 1.0
STDs:syphilis            1.0
STDs:pelvic inflammatory disease 0.0
STDs:genital herpes       0.0
STDs:molluscum contagiosum 0.0
STDs:AIDS                0.0
STDs:HIV                 1.0
STDs:Hepatitis B          0.0
STDs:HPV                 0.0
STDs: Number of diagnosis 1.0
STDs: Time since first diagnosis 12.0
STDs: Time since last diagnosis 12.0
Dx:Cancer                1.0
Dx:CIN                  1.0
Dx:HPV                  1.0
Dx                      1.0
Hinselmann               1.0
Schiller                 1.0
Cytology                 1.0
dtype: float64

```

```
X_train[numerical_columns].describe()
```

	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	...	STDs: Number of diagnosis	STDs: Time since first diagnosis	STDs: Time since last diagnosis	Dx:Cancer	Dx:CIN	Dx:HPV
664.000000	664.000000	664.000000	664.000000	...	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000	664
0.475850	0.691265	1.916573	0.096386	...	0.096386	1.487952	1.453313	0.019578	0.012048	0.022590	0
2.332357	0.462320	3.499791	0.295342	...	0.319859	2.414731	2.312192	0.138650	0.109183	0.148706	0
0.000000	0.000000	0.000000	0.000000	...	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0
0.000000	0.000000	0.000000	0.000000	...	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0
0.000000	1.000000	0.250000	0.000000	...	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0
0.000000	1.000000	2.000000	0.000000	...	0.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0
37.000000	1.000000	30.000000	1.000000	...	3.000000	22.000000	22.000000	1.000000	1.000000	1.000000	1

## ▼ Feature Scaling

```
X_train.describe()
```

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD
count	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000	664.000000
mean	27.060241	2.484940	17.025602	2.25000	0.146084	1.258027	0.475850	0.691265	1.916573	0.096386
std	7.905894	1.318078	2.782884	1.39299	0.353457	4.218211	2.332357	0.462320	3.499791	0.295342
min	14.000000	1.000000	10.000000	0.00000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	21.000000	2.000000	15.000000	1.00000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	26.000000	2.000000	17.000000	2.00000	0.000000	0.000000	0.000000	1.000000	0.250000	0.000000
75%	32.000000	3.000000	18.000000	3.00000	0.000000	0.000000	0.000000	1.000000	2.000000	0.000000
max	70.000000	10.000000	32.000000	8.00000	1.000000	37.000000	37.000000	1.000000	30.000000	1.000000

8 rows × 35 columns

```
cols = X_train.columns

from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
X_train= scaler.fit_transform(X_train)
X_test= scaler.transform(X_test)

X_train= pd.DataFrame (X_train, columns=[cols])
```

```
X_train.describe()
```

8 rows × 35 columns

## ▼ Model Training

```
from sklearn.linear_model import LogisticRegression

# Instantiate the logistic regression model
logreg = LogisticRegression(solver='liblinear', random_state=42) # Adjust solver and random_state as needed

# Fit the model
logreg.fit(X_train, y_train)
```

## ▼ Predict results

```
0.98030213, 0.98187561, 0.97602653, 0.40791458, 0.98165476,
0.95004821, 0.97979568, 0.97847942, 0.97929069, 0.93925603,
0.98250452, 0.9709148 , 0.98038387, 0.69087378, 0.98028688,
0.97448159, 0.9803807 , 0.98292581, 0.97990914, 0.98114977,
0.97404099, 0.98294775, 0.96048621, 0.98351494, 0.98593851,
0.97606419, 0.9565455 , 0.98057819, 0.97549488, 0.98791546,
0.98474134, 0.97960026, 0.35920079, 0.9704203 , 0.94747769,
0.97948509, 0.97877902, 0.40118625, 0.96721231, 0.98209548,
0.98650188, 0.97285332, 0.98067339, 0.97352784, 0.98302117,
0.98577423, 0.96692052, 0.98365639, 0.98816839, 0.48103242,
0.98057753, 0.28532274, 0.98055511, 0.98300654, 0.98142108,
0.97883932, 0.97356948, 0.97790029, 0.98877925, 0.95339285,
0.98452563, 0.98644564, 0.97396277, 0.98040971, 0.98585397,
0.97956192, 0.2697327 , 0.97700375, 0.98761238, 0.98325658,
0.98468998, 0.98058107, 0.9841359 , 0.97569208, 0.98068851,
0.97578687, 0.9608319 , 0.98490733, 0.64426007, 0.98393871,
0.64400598, 0.98369889, 0.31542777, 0.97849007, 0.96344689,
0.43346647, 0.98230825, 0.97971349, 0.98054514, 0.98185996,
0.98005718, 0.9780225 , 0.87323158, 0.98384355, 0.98372832,
0.97934297, 0.96756457, 0.25295199, 0.97736216, 0.97063763,
0.98701769, 0.9792027 , 0.9850321 , 0.98460897, 0.73267532,
0.9839528 , 0.98210336, 0.34175938, 0.96998096, 0.97882308,
0.98211513, 0.99030635, 0.98206261, 0.98095815, 0.98053311,
0.55083532, 0.97598307, 0.97103627, 0.17077463, 0.98301245,
0.99086979, 0.97994397, 0.98990966, 0.98011905, 0.97113264,
0.95499399, 0.97033208, 0.98637817, 0.97965374, 0.23696323,
0.97727537, 0.9802864 , 0.98666697, 0.98364211, 0.4234108 ,
0.96269302, 0.97685968, 0.9750471 , 0.9785283 , 0.98229493,
0.9892733 , 0.95081616, 0.97757596, 0.98498087, 0.98062242,
0.98132382, 0.97882584, 0.98487212, 0.97353172, 0.97581305,
0.9792013 , 0.96751914, 0.98552693, 0.97980147, 0.96496642,
0.97813663])
```

```
logreg.predict_proba(X_test)[:,1]
```

```
array([0.02832186, 0.01378096, 0.02451266, 0.02027214, 0.01632066,
       0.02303176, 0.01549106, 0.01637031, 0.01568066, 0.02300414,
       0.01969787, 0.01812439, 0.02397347, 0.59208542, 0.01834524,
       0.04995179, 0.02020432, 0.02152058, 0.02070931, 0.06074397,
       0.01749548, 0.0290852 , 0.01961613, 0.30912622, 0.01971312,
       0.02551841, 0.0196193 , 0.01707419, 0.02009086, 0.01885023,
       0.02595901, 0.01705225, 0.03951379, 0.01648506, 0.01406149,
       0.02393581, 0.0434545 , 0.01942181, 0.02450512, 0.01208454,
       0.01525866, 0.02039974, 0.64079921, 0.0295797 , 0.05252231,
       0.02051491, 0.02122098, 0.59881375, 0.03278769, 0.01790452,
       0.01349812, 0.02714668, 0.01932661, 0.02647216, 0.01697883,
       0.01422577, 0.03307948, 0.01634361, 0.01183161, 0.51896758,
       0.01942247, 0.71467726, 0.01944489, 0.01699346, 0.01857892,
       0.02116068, 0.02643052, 0.02209971, 0.01122075, 0.04660715,
       0.01547437, 0.01355436, 0.02603723, 0.01959029, 0.01414603,
       0.02043808, 0.7302673 , 0.02299625, 0.01238762, 0.01674342,
       0.01531002, 0.01941893, 0.0158641 , 0.02430792, 0.01931149,
       0.02421313, 0.0391681 , 0.01509267, 0.35573993, 0.01606129,
       0.35599402, 0.01630111, 0.68457223, 0.02150993, 0.03655311,
       0.56653353, 0.01769175, 0.02028651, 0.01945486, 0.01814004,
       0.01994282, 0.0219775 , 0.12676842, 0.01615645, 0.01627168,
       0.02065703, 0.03243543, 0.74704801, 0.02263784, 0.02936237,
       0.01298231, 0.0207973 , 0.0149679 , 0.01539103, 0.26732468,
       0.0160472 , 0.01789664, 0.65824062, 0.03001904, 0.02117692,
       0.01788487, 0.00969365, 0.01793739, 0.01904185, 0.01946689,
       0.44916468, 0.02401693, 0.02896373, 0.82922537, 0.01698755,
       0.00913021, 0.02005603, 0.01009034, 0.01988095, 0.02886736,
       0.04500601, 0.02966792, 0.01362183, 0.02034626, 0.76303677,
       0.02272463, 0.0197136 , 0.01333303, 0.01635789, 0.5765892 ,
       0.03730698, 0.02314032, 0.0249529 , 0.0214717 , 0.01770507,
       0.0107267 , 0.04918384, 0.02242404, 0.01501913, 0.01937758,
       0.01867618, 0.02117416, 0.01512788, 0.02646828, 0.02418695,
       0.0207987 , 0.03248086, 0.01447307, 0.02019853, 0.03503358,
       0.02186337])
```

## ✓ Check accuracy score

```
from sklearn.metrics import accuracy_score
print('Model accuracy score: {:.4f}'.format(accuracy_score(y_test, y_pred_test)))
```

```
Model accuracy score: 0.9639
```

Compare the train-set and test-set accuracy Now, I will compare the train-set and test-set accuracy to check for overfitting.

```
print('Training accuracy score: {:.4f}'.format(accuracy_score(y_train, y_pred_train)))
```

Training accuracy score: 0.9623

## Check for overfitting and underfitting

```
#print the scores on training and test set
print('Training set score: {:.4f}'.format(logreg.score(X_train, y_train)))
print('Test set score: {:.4f}'.format(logreg.score(X_test, y_test)))
```

Training set score: 0.9623  
Test set score: 0.9639

```
#fit the Logisitic Regression model with C=100
#instantiate the model
logreg100 = LogisticRegression (C=100, solver='liblinear', random_state=0)
#fit the model
logreg100.fit(X_train, y_train)
```

```
LogisticRegression(C=100, random_state=0, solver='liblinear')  
  
#print the scores on training and test set
```

Training set score: 0.9654

```
#check class distribution in test set
```

```
Biopsy
0    155
1     11
Name: count, dtype: int64
```

```
#check null accuracy score
null_accuracy= (155/(155+11))
print('Null accuracy score: {0:.4f}'. format(null_accuracy))

Null accuracy score: 0.9337
```

## ▼ Confusion Matrix


```
#Print the Confusion Matrix and slice it into four pieces
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred_test)
print('Confusion matrix\n\n', cm)
print('\nTrue Positives (TP)=', cm[0,0])
print('\nTrue Negatives (TN)=', cm[1,1])
print('\nFalse Positives (FP)=', cm[0,1])
print('\nFalse Negatives (FN)=', cm[1,0])

Confusion matrix

[[151  4]
 [ 2  9]]

True Positives (TP)= 151
True Negatives (TN)= 9
False Positives (FP)= 4
False Negatives (FN)= 2
```

```
#visualize confusion matrix with seaborn heatmap
cm_matrix=pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                       index=['Predict Positive:1', 'Predict Negative:0'])
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```



The confusion matrix represents the performance of a classification model. It shows the number of true positives (correctly predicted positive instances), true negatives (correctly predicted negative instances), false positives (incorrectly predicted positive instances), and false negatives (incorrectly predicted negative instances). The model has high true positives, indicating good performance in identifying the positive class. However, there are more false positives than false negatives, suggesting potential bias. Further improvements could enhance precision and recall.

## Classification Metrics

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_test))

      precision    recall  f1-score   support

       0       0.99    0.97    0.98     155
       1       0.69    0.82    0.75      11

  accuracy                           0.96    166
  macro avg       0.84    0.90    0.87    166
weighted avg       0.97    0.96    0.97    166
```

Classification accuracy

```
TP= cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]

#print classification accuracy
classification_accuracy=(TP+TN) / float(TP + TN+FP + FN)
print('Classification accuracy: {0:0.4f}'.format(classification_accuracy))

Classification accuracy: 0.9639
```

Classification error

```
#print classification error
classification_error= (FP+ FN) / float (TP+ TN+FP+ FN)
print('Classification error: {0:0.4f}'.format(classification_error))

Classification error: 0.0361
```

Precision

```
#print precision score
precision=TP/float(TP+FP)
print('Precision {0:0.4f}'.format(precision))

Precision 0.9742
```

Recall

```
recall=TP/float (TP + FN)
print('Recall or Sensitivity: {0:0.4f}'.format(recall))

Recall or Sensitivity: 0.9869
```

True Positive Rate

```
true_positive_rate= TP / float(TP + FN)
print('True Positive Rate: {0:0.4f}'.format(true_positive_rate))

True Positive Rate: 0.9869
```

False Positive Rate

```
false_positive_rate= FP / float (FP + TN)
print('False Positive Rate: {0:0.4f}'.format(false_positive_rate))

False Positive Rate: 0.3077
```

## Specificity

```
specificity=TN/ (TN+FP)
print('Specificity {0:.4f}'.format(specificity))
```

Specificity 0.6923

## Adjusting the threshold level

```
#print the first 18 predicted probabilities of two classes 0 and 1
y_pred_prob = logreg.predict_proba(X_test)[8:10]
y_pred_prob

array([[0.98431934, 0.01568066],
       [0.97699586, 0.02300414]])
```

---

```
#store the probabilities in dataframe
y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of No cancer (0)', 'Prob of cancer (1)'])
y_pred_prob_df
```

	Prob of No cancer (0)	Prob of cancer (1)
0	0.984319	0.015681
1	0.976996	0.023004

Next steps: [View recommended plots](#)

```
logreg.predict_proba (X_test) [0:10, 1]

array([0.02832186, 0.01378096, 0.02451266, 0.02027214, 0.01632066,
       0.02303176, 0.01549106, 0.01637031, 0.01568066, 0.02300414])

y_pred=logreg.predict_proba(X_test)[:, 1]

# Set font size
plt.rcParams['font.size'] = 12


# Plot histogram with 10 bins
plt.hist(y_pred, bins=10)

# Set title
plt.title('Histogram of Predicted Probabilities of Cancer')

# Set x-axis limit
plt.xlim(0, 1)

# Set axis labels
plt.xlabel('Predicted Probabilities of Cancer')
plt.ylabel('Frequency')

# Show plot
plt.show()
```



The majority of predictions are concentrated around a low probability range (0-0.2), indicating that most individuals in this dataset have a low predicted risk of cancer. There are no predictions showing a moderate to high probability, which suggests that, according to this model, the studied population is at low risk for cancer.

Lower the threshold

```
from sklearn.preprocessing import binarize
from sklearn.metrics import confusion_matrix, accuracy_score

for i in range(1, 5):
    cm1 = 0
    y_pred1 = logreg.predict_proba(X_test)[:, 1].reshape(-1, 1)
    y_pred2 = binarize(y_pred1, threshold=i / 10) # Correct usage of binarize
    cm1 = confusion_matrix(y_test, y_pred2)
    print('With', i / 10, 'threshold the Confusion Matrix is:\n\n', cm1, '\n\n',
          'With', cm1[0, 0] + cm1[1, 1], 'correct predictions\n\n',
          cm1[0, 1], 'Type I errors (False Positives)\n\n',
          cm1[1, 0], 'Type II errors (False Negatives)\n\n',
          'Accuracy score:', accuracy_score(y_test, y_pred2), '\n\n',
          'Sensitivity:', cm1[1, 1] / (float(cm1[1, 1] + cm1[1, 0])), '\n\n',
          'Specificity:', cm1[0, 0] / (float(cm1[0, 0] + cm1[0, 1])), '\n\n',
          '=====', '\n\n')
```

With 0.2 threshold the Confusion Matrix is:

```
[[148  7]
 [ 0 11]]
```

With 159 correct predictions

```
6 Type I errors (False Positives)
```

```
0 Type II errors (False Negatives)
```

```
Accuracy score: 0.963855421686747
```

```
Sensitivity: 1.0
```

```
Specificity: 0.9612903225806452
```

```
=====
```

With 0.4 threshold the Confusion Matrix is:

```
[[150  5]
 [ 2  9]]
```

With 159 correct predictions

```
5 Type I errors (False Positives)
```

```
2 Type II errors (False Negatives)
```

```
Accuracy score: 0.9578313253012049
```

```
Sensitivity: 0.81818181818182
```

```
Specificity: 0.967741935483871
```


```
=====
```

## ▼ ROC-AUC

```
from sklearn.metrics import roc_curve, auc

# Calculate ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(6, 4))
plt.plot(fpr, tpr, linewidth=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.title('ROC curve for Cervical Cancer classifier')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.legend(loc='lower right')
plt.show()
```



The graph is a Receiver Operating Characteristic (ROC) curve for a Cervical Cancer classifier. The ROC curve is close to the top left corner, indicating high sensitivity and low false positive rate. With an Area Under Curve (AUC) of 0.99, the classifier demonstrates excellent performance, suggesting it is highly effective and reliable for predicting cervical cancer.

```
# Make predictions on the test set
y_pred = logreg.predict(X_test)

# Output predictions for each individual
for i, prediction in enumerate(y_pred):
    if prediction == 1:
        print(f"Index {i} prediction: will have cancer")
    else:
```