

Linear Regression Analysis: Automobile Datasets

Marquez, Keith Leigh Zhen R.

▼ Data Wrangling:

```
pip install ucimlrepo

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)

from ucimlrepo import fetch_ucirepo

# fetch dataset
automobile = fetch_ucirepo(id=10)

# data (as pandas dataframes)
X = automobile.data.features
y = automobile.data.targets

# metadata
print(automobile.metadata)

# variable information
print(automobile.variables)

[{"uci_id": 10, "name": "Automobile", "repository_url": "https://archive.ics.uci.edu/dataset/10/automobile", "data_url": "https://archive.ics.uci.edu/ml/machine-learning-databases/00330/automobile.names"}, {"name": "price", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "highway-mpg", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "city-mpg", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "peak-rpm", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "horsepower", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "compression-ratio", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "stroke", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "bore", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "fuel-system", "role": "Feature", "type": "Categorical", "demographic": "None"}, {"name": "engine-size", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "num-of-cylinders", "role": "Feature", "type": "Integer", "demographic": "None"}, {"name": "engine-type", "role": "Feature", "type": "Categorical", "demographic": "None"}, {"name": "curb-weight", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "height", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "width", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "length", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "wheel-base", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "engine-location", "role": "Feature", "type": "Binary", "demographic": "None"}, {"name": "drive-wheels", "role": "Feature", "type": "Categorical", "demographic": "None"}, {"name": "body-style", "role": "Feature", "type": "Categorical", "demographic": "None"}, {"name": "num-of-doors", "role": "Feature", "type": "Integer", "demographic": "None"}, {"name": "aspiration", "role": "Feature", "type": "Binary", "demographic": "None"}, {"name": "fuel-type", "role": "Feature", "type": "Binary", "demographic": "None"}, {"name": "make", "role": "Feature", "type": "Categorical", "demographic": "None"}, {"name": "normalized-losses", "role": "Feature", "type": "Continuous", "demographic": "None"}, {"name": "symboling", "role": "Target", "type": "Integer", "demographic": "None"}]

[{"description": "continuous from 5118 to 45400", "units": "None", "missing_values": "yes"}, {"description": "continuous from 16 to 54", "units": "None", "missing_values": "no"}, {"description": "continuous from 13 to 49", "units": "None", "missing_values": "no"}, {"description": "continuous from 4150 to 6600", "units": "None", "missing_values": "yes"}, {"description": "continuous from 48 to 288", "units": "None", "missing_values": "yes"}, {"description": "continuous from 7 to 23", "units": "None", "missing_values": "no"}, {"description": "continuous from 2.07 to 4.17", "units": "None", "missing_values": "yes"}, {"description": "continuous from 2.54 to 3.94", "units": "None", "missing_values": "yes"}, {"description": "1bbl, 2bbl, 4bbl, id1, mfi, mpfi, spdi, spfi", "units": "None", "missing_values": "no"}, {"description": "continuous from 61 to 326", "units": "None", "missing_values": "no"}, {"description": "eight, five, four, six, three, twelve, two", "units": "None", "missing_values": "no"}, {"description": "dohc, dohcvt, l, ohc, ohcf, ohcv, rotor", "units": "None", "missing_values": "no"}, {"description": "continuous from 1488 to 4066", "units": "None", "missing_values": "no"}, {"description": "continuous from 47.8 to 59.8", "units": "None", "missing_values": "no"}, {"description": "continuous from 60.3 to 72.3", "units": "None", "missing_values": "no"}, {"description": "continuous from 141.1 to 208.1", "units": "None", "missing_values": "no"}, {"description": "continuous from 86.6 120.9", "units": "None", "missing_values": "no"}, {"description": "front, rear", "units": "None", "missing_values": "no"}, {"description": "4wd, fwd, rwd", "units": "None", "missing_values": "no"}]
```

```

19      hardtop, wagon, sedan, hatchback, convertible  None      no
20                  four, two   None      yes
21                  std, turbo  None      no
22                  diesel, gas  None      no
23 alfa-romero, audi, bmw, chevrolet, dodge, hond...  None      no
24                  continuous from 65 to 256  None      yes
25                  -3, -2, -1, 0, 1, 2, 3  None      no

```

```

import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```

auto = pd.concat([X,y], axis=1)
auto

```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	stroke	bore	fuel-system	engine-size	...	wheel-base	engine-location	drive-wheels	body-style	c
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	130	...	88.6	front	rwd	convertible	
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	130	...	88.6	front	rwd	convertible	
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	mpfi	152	...	94.5	front	rwd	hatchback	
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	mpfi	109	...	99.8	front	fwd	sedan	
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	mpfi	136	...	99.4	front	4wd	sedan	
...	
200	16845.0	28	23	5400.0	114.0	9.5	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	
201	19045.0	25	19	5300.0	160.0	8.7	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	
202	21485.0	23	18	5500.0	134.0	8.8	2.87	3.58	mpfi	173	...	109.1	front	rwd	sedan	
203	22470.0	27	26	4800.0	106.0	23.0	3.40	3.01	idi	145	...	109.1	front	rwd	sedan	
204	22625.0	25	19	5400.0	114.0	9.5	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	

205 rows × 26 columns

```
auto.dtypes
```

price	float64
highway-mpg	int64
city-mpg	int64
peak-rpm	float64
horsepower	float64
compression-ratio	float64
stroke	float64
bore	float64
fuel-system	object
engine-size	int64
num-of-cylinders	int64
engine-type	object
curb-weight	int64
height	float64
width	float64
length	float64
wheel-base	float64
engine-location	object
drive-wheels	object
body-style	object
num-of-doors	float64
aspiration	object
fuel-type	object
make	object
normalized-losses	float64
symboling	int64
dtype: object	

Identify missing value

```
auto_null=auto.isnull().sum()
auto_null

  price          4
highway-mpg       0
city-mpg         0
peak-rpm         2
horsepower       2
compression-ratio 0
stroke          4
bore             4
fuel-system      0
engine-size      0
num-of-cylinders 0
engine-type      0
curb-weight      0
height           0
width            0
length           0
wheel-base        0
engine-location   0
drive-wheels     0
body-style        0
num-of-doors      2
aspiration        0
fuel-type         0
make              0
normalized-losses 41
symboling         0
dtype: int64
```

Checking for duplicates

```
# Check for duplicates
duplicate_rows = auto.duplicated()

# Count the number of duplicate rows
num_duplicates = duplicate_rows.sum()
print("Number of duplicate rows:", num_duplicates)

Number of duplicate rows: 0
```

Summary of the data

```
auto.describe()
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	stroke	bore	engine-size	num-of-cylinders	cu
count	201.000000	205.000000	205.000000	203.000000	203.000000	205.000000	201.000000	201.000000	205.000000	205.000000	205.000000
mean	13207.129353	30.751220	25.219512	5125.369458	104.256158	10.142537	3.255423	3.329751	126.907317	4.380488	2555.5651
std	7947.066342	6.886443	6.542142	479.334560	39.714369	3.972040	0.316717	0.273539	41.642693	1.080854	520.680:
min	5118.000000	16.000000	13.000000	4150.000000	48.000000	7.000000	2.070000	2.540000	61.000000	2.000000	1488.000
25%	7775.000000	25.000000	19.000000	4800.000000	70.000000	8.600000	3.110000	3.150000	97.000000	4.000000	2145.000
50%	10295.000000	30.000000	24.000000	5200.000000	95.000000	9.000000	3.290000	3.310000	120.000000	4.000000	2414.000
75%	16500.000000	34.000000	30.000000	5500.000000	116.000000	9.400000	3.410000	3.590000	141.000000	4.000000	2935.000
max	45400.000000	54.000000	49.000000	6600.000000	288.000000	23.000000	4.170000	3.940000	326.000000	12.000000	4066.000

Fill in missing data

```
auto['price'].fillna(auto['price'].median(), inplace=True)
auto['peak-rpm'].fillna(auto['peak-rpm'].median(), inplace=True)
auto['horsepower'].fillna(auto['horsepower'].median(), inplace=True)
auto['stroke'].fillna(auto['stroke'].median(), inplace=True)
auto['bore'].fillna(auto['bore'].median(), inplace=True)
auto['num-of-doors'].fillna(auto['num-of-doors'].mode()[0], inplace=True)
auto['normalized-losses'].fillna(auto['normalized-losses'].median(), inplace=True)
auto
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	stroke	bore	fuel-system	engine-size	...	wheel-base	engine-location	drive-wheels	body-style	c
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	130	...	88.6	front	rwd	convertible	
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	130	...	88.6	front	rwd	convertible	
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	mpfi	152	...	94.5	front	rwd	hatchback	
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	mpfi	109	...	99.8	front	fwd	sedan	
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	mpfi	136	...	99.4	front	4wd	sedan	
...
200	16845.0	28	23	5400.0	114.0	9.5	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	
201	19045.0	25	19	5300.0	160.0	8.7	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	
202	21485.0	23	18	5500.0	134.0	8.8	2.87	3.58	mpfi	173	...	109.1	front	rwd	sedan	
203	22470.0	27	26	4800.0	106.0	23.0	3.40	3.01	idi	145	...	109.1	front	rwd	sedan	
204	22625.0	25	19	5400.0	114.0	9.5	3.15	3.78	mpfi	141	...	109.1	front	rwd	sedan	

205 rows × 26 columns

Checking for null

auto.isnull().sum()

```

price          0
highway-mpg    0
city-mpg        0
peak-rpm        0
horsepower      0
compression-ratio 0
stroke          0
bore            0
fuel-system     0
engine-size     0
num-of-cylinders 0
engine-type     0
curb-weight     0
height          0
width           0
length          0
wheel-base      0
engine-location 0
drive-wheels    0
body-style       0
num-of-doors    0
aspiration      0
fuel-type        0
make            0
normalized-losses 0
symboling       0
dtype: int64

```

Convert/change data types

```

auto['num-of-doors'] = auto['num-of-doors'].astype(int)
auto.dtypes

```

```

price          float64
highway-mpg    int64
city-mpg        int64
peak-rpm        float64
horsepower      float64
compression-ratio float64
stroke          float64
bore            float64
fuel-system     object
engine-size     int64
num-of-cylinders int64
engine-type     object

```

```

curb-weight      int64
height          float64
width           float64
length          float64
wheel-base      float64
engine-location object
drive-wheels    object
body-style      object
num-of-doors    int64
aspiration      object
fuel-type       object
make            object
normalized-losses float64
symboling       int64
dtype: object

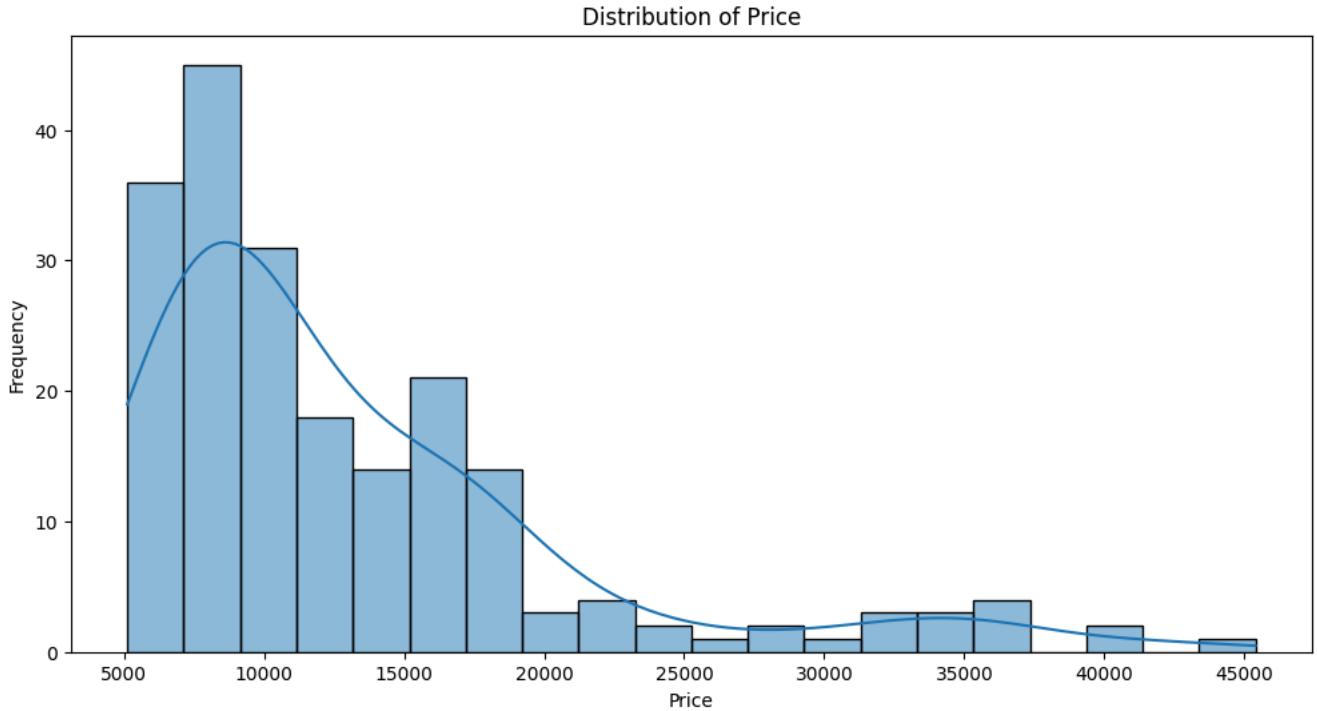
```

✓ EDA (exploratory data analysis):

```

plt.figure(figsize=(12, 6))
sns.histplot(auto['price'], bins=20, kde=True)
plt.title('Distribution of Price')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()

```



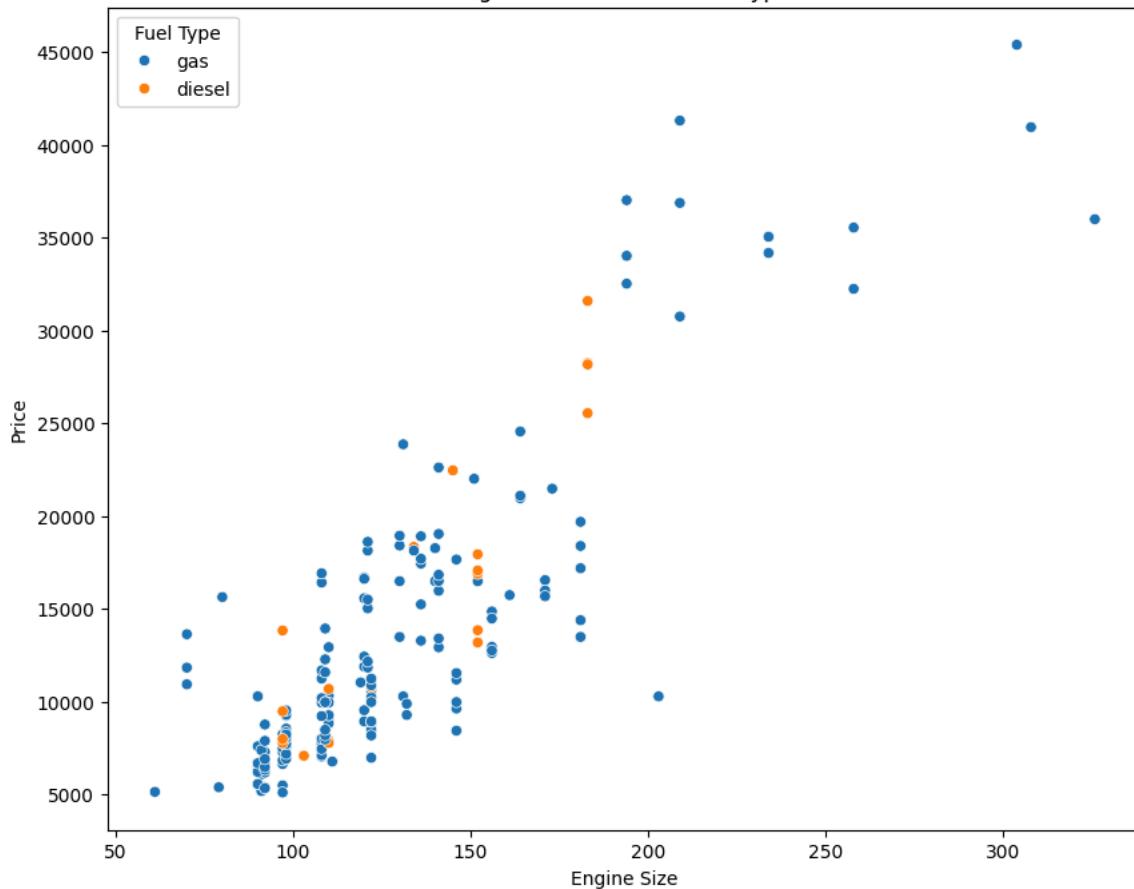
The graph shows a Distribution of Price with an overlaid line graph. It reveals that the majority of data points cluster around the 10,000 mark, indicating it's the most common price. Beyond this point, there's a decrease in frequency, suggesting higher-priced items are less common. The downward trend of the line graph reinforces this, indicating fewer items as prices rise. This implies a market where lower-priced items dominate and consumers may be sensitive to price increases.

```

plt.figure(figsize=(10, 8))
sns.scatterplot(data=auto, x='engine-size', y='price', hue='fuel-type')
plt.title('Engine Size vs Price (Fuel Type)')
plt.xlabel('Engine Size')
plt.ylabel('Price')
plt.legend(title='Fuel Type')
plt.show()

```

Engine Size vs Price (Fuel Type)



The graph displays a scatter plot comparing engine size to price, with distinctions made for gas and diesel fuel types. Gas vehicles (depicted in blue) outnumber diesel vehicles (depicted in orange) and generally offer a wider range of engine sizes and prices, making them more affordable, especially for smaller engines. Diesel vehicles tend to have mid-range engine sizes and higher prices. This suggests that consumers seeking budget-friendly options may prefer gas vehicles with smaller engines, while diesel vehicles, despite their higher cost, may offer benefits justifying the expense, particularly for mid-range engine sizes. Overall, the trend indicates that larger engine sizes typically correlate with higher prices for both fuel types, though there are exceptions, especially among gas vehicles. This implies that factors beyond engine size, such as brand, features, or market demand, may also influence pricing.

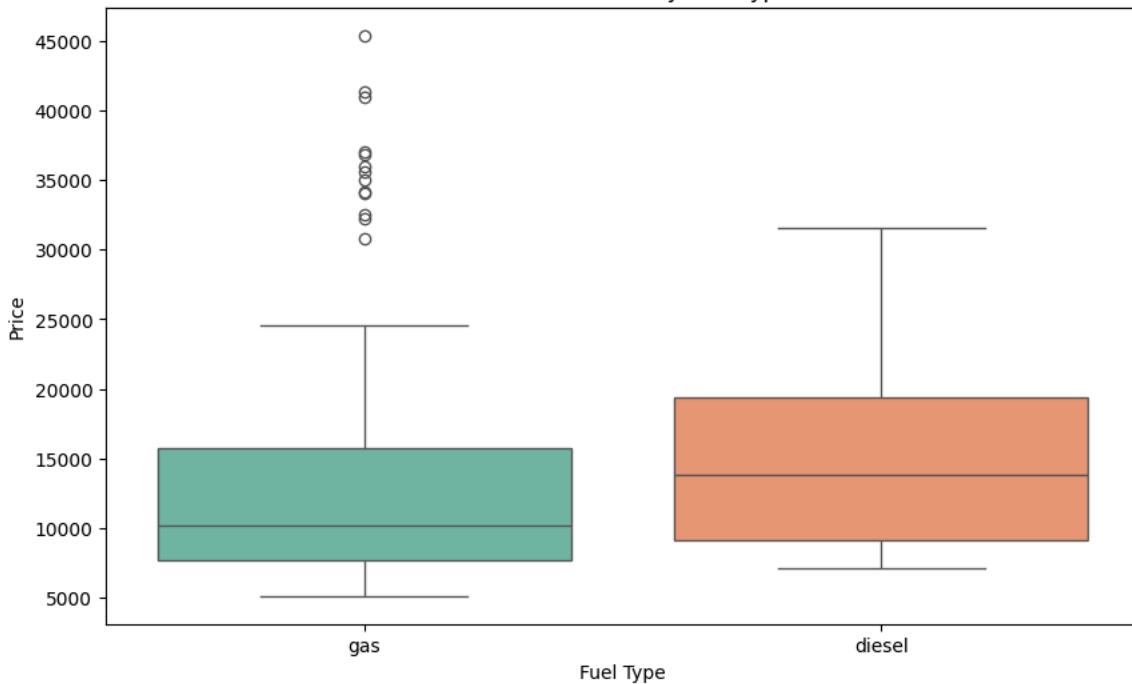
```
plt.figure(figsize=(10, 6))
sns.boxplot(x='fuel-type', y='price', data=auto, palette='Set2')
plt.title('Price Distribution by Fuel Type')
plt.xlabel('Fuel Type')
plt.ylabel('Price')
plt.show()
```

```
<ipython-input-20-5742f86d257d>:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

```
sns.boxplot(x='fuel-type', y='price', data=auto, palette='Set2')
```

Price Distribution by Fuel Type



The graph compares vehicle prices by fuel type, showing gas vehicles have a median price around 10,000, with an interquartile range (IQR) of 7,500 to 12,500 and outliers reaching 45,000. Diesel vehicles have a higher median price of approximately 17,500, with a broader IQR (15,000 to 20,000) indicating greater price variance. Overall, diesel cars are generally pricier upfront, while gas cars exhibit more variability, occasionally surpassing diesel prices due to specific models or features. Gas vehicles include many high-priced outliers, while diesel prices are more consistently higher.

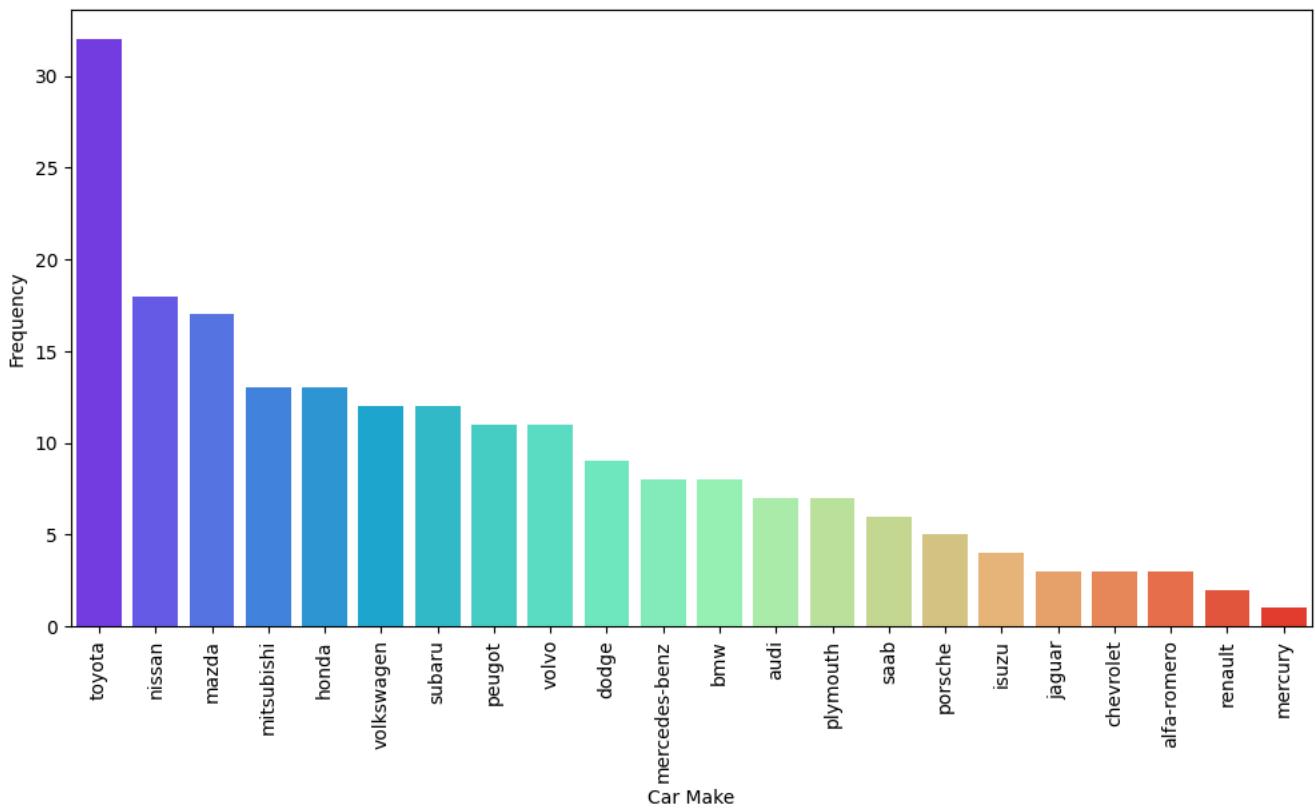
```
plt.figure(figsize=(12, 6))
sns.countplot(x='make', data=auto, order=auto['make'].value_counts().index, palette='rainbow')
plt.xticks(rotation=90)
plt.title('Distribution of Car Makes')
plt.xlabel('Car Make')
plt.ylabel('Frequency')
plt.show()
```

```
<ipython-input-19-6b7cb4b082c2>:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend`

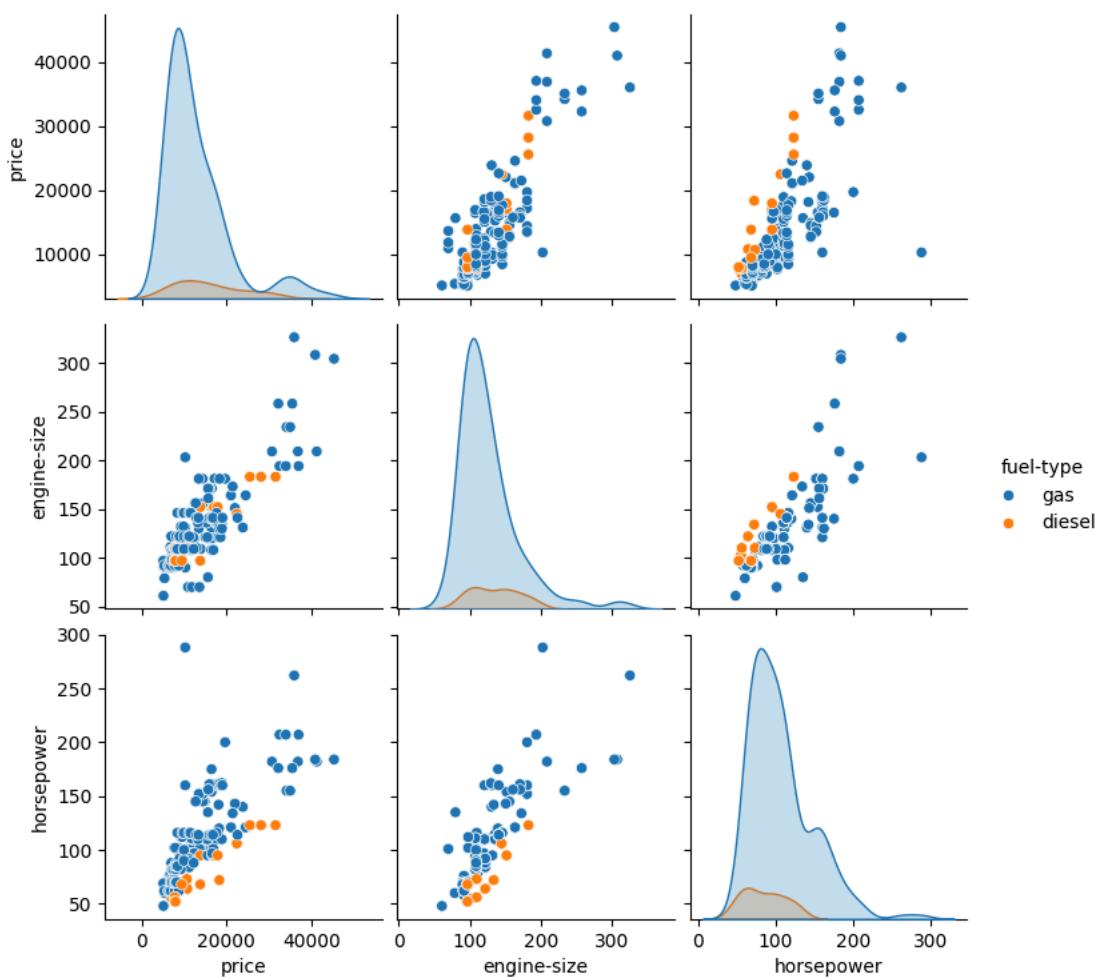
```
sns.countplot(x='make', data=auto, order=auto['make'].value_counts().index, palette='rainbow')
```

Distribution of Car Makes



The graph illustrates the distribution of car makes, revealing Toyota's dominance with over 30% frequency, followed by Nissan and Mazda. Luxury brands like Mercedes-Benz, BMW, and Audi hold moderate frequencies, while less common makes such as Jaguar, Alfa Romeo, Renault, and Mercury have limited popularity. This indicates Toyota's significant lead in the market, with other brands competing in a diverse but competitive landscape.

```
sns.pairplot(auto[['price', 'engine-size', 'horsepower', 'fuel-type']], hue='fuel-type')
plt.show()
```



This shows that cars with larger engines and higher horsepower tend to command higher prices, particularly evident in diesel cars. While there's a clear positive correlation between horsepower and price, the link between engine size and horsepower is less distinct. Overall, the visual representations highlight the trend of higher specifications correlating with higher prices, emphasizing the importance of considering fuel type when analyzing these relationships.

Linear Regression Analysis

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

Selecting predictor variables and target variable

```
predictors = ['engine-size', 'horsepower', 'curb-weight', 'highway-mpg']
target = 'price'
```

Splitting the dataset into training and testing sets

```
X = auto[predictors]
y = auto[target]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Train the linear regression model

```
model = LinearRegression()
model.fit(X_train, y_train)
```

```
└─ LinearRegression  
    └─ LinearRegression()
```

Make predictions on the test set

```
y_pred = model.predict(X_test)
```

Evaluate the model

```
mse = metrics.mean_squared_error(y_test, y_pred)
```